Type Systems as Macros

Stephen Chang   Alex Knauth   Ben Greenman

PLT @ Northeastern University
{stchang,alexknauth,types}@ccs.neu.edu

Abstract

We present TURNSTILE, a metalanguage for creating typed embedded languages. To implement the type system, programmers write type checking rules resembling traditional judgment syntax. To implement the semantics, they incorporate elaborations into these rules. TURNSTILE critically depends on the idea of linguistic reuse. It exploits a *macro system* in a novel way to *simultaneously* type check and rewrite a surface program into a target language. Reusing a macro system also yields modular implementations whose rules may be mixed and matched to create other languages. Combined with typical compiler and runtime reuse, TURNSTILE produces performant typed languages with little effort.

1. Typed Embedded Languages

As Paul Hudak asserted, "we really don't want to build a programming language from scratch ... better, let's inherit the infrastructure of some other language" [24]. Unsurprisingly, many modern languages support the creation of such embedded languages [4, 19, 21, 22, 25, 26, 28, 40, 42, 45].

Programmers who wish to create *typed* embedded languages, however, have more limited options. Such languages typically reuse their host's type system but, as a prominent project [44] recently remarked, this "*confines* them to ⟨that⟩ type system." Also, reusing a type system may not create proper abstractions, e.g., type errors may be reported in host language terms. At the other extreme, a programmer can implement a type system from scratch [41], expending considerable effort and passing up many of the reuse benefits that embedding a language promises in the first place.

We present an alternative approach to implementing typed embedded languages. Rather than reuse a type system, we embed a type system in a host's *macro system*. In other words, type checking is computed as part of macro expansion. Such an embedding fits naturally since a typical type checking algorithm traverses a surface program, synthesizes information from it, and uses this information to rewrite the program, if it satisfies certain conditions, into a target language. *This kind of algorithm exactly matches the ideal use case for macros.* From this perspective, a type checker resembles a special instance of a macro system and our approach exploits synergies resulting from this insight.

With our macro-based approach, programmers may implement a wide range of type rules, yet they need not create a type system from scratch as they may reuse many parts of a macro system for type checking. Indeed, programmers need only supply their desired type rules in an intuitive mathematical form. Creating type systems with macros also fosters robust linguistic abstractions, e.g., they report type errors with surface language terms. Finally, our approach produces naturally modular type systems that dually serve as libraries of mixable and matchable type rules, enabling further linguistic reuse [29]. When combined with the typical reuse of the runtime that embedded languages enjoy, our approach inherits the performance of its host and thus produces practical typed languages with little effort.

We use Racket [13, 16], a Lisp and Scheme descendant, as our host language since Lisps are already a popular platform for creating embedded languages [18, 21]. Racket's macro system in particular continues to improve on its predecessors [15] and has even influenced macro system design in modern non-Lisp languages [7, 9, 11, 47]. Thus programmers have created Racket-embedded languages for accomplishing a variety of tasks such as book publishing [8], program synthesis [43], and writing secure shell scripts [33].

The first part of the paper (§2-3) demonstrates a connection between type rules and macros by reusing Racket's macro infrastructure for type checking in the creation of a typed embedded language. The second part (§4) introduces TURNSTILE, a metalanguage that abstracts the insights and techniques from the first part into convenient linguistic constructs. The third part (§5-7) shows that our approach both accommodates a variety of type systems and scales to realistic combinations of type system features. We demonstrate the former by implementing fifteen core languages ranging from simply-typed to F_ω, and the latter with the creation of a full-sized ML-like functional language that also supports Haskell-style type classes.

2. Creating Embedded Languages in Racket

This section summarizes the creation of embedded languages with Racket. Racket is not a single language but rather an ecosystem with which to create languages [13]. Racket code is organized into modules, e.g. LAM:[1]

```
#lang racket                                    LAM
(define-m (lm x e)  (λ (x) e))
(provide lm)
```

A `#lang racket` declaration allows LAM to use forms and functions from the main Racket language. LAM defines and exports one macro, `lm`, denoting single-argument functions. A Racket macro[2] consumes and produces *syntax object* data structures. The `lm` macro specifies its usage shape with input pattern `(lm x e)` (in yellow to help readability[3]), which binds *pattern variables* `x` and `e` to subpieces of the input, the parameter and body, respectively. The output syntax `(λ (x) e)` (gray denotes syntax object construction) references these pattern variables (λ is Racket's λ).

A module serves multiple roles in the Racket ecosystem. Running LAM as a *program* produces no result since it consists of only a macro definition. But LAM is also a *language*:

```
#lang lam                                  LAM-PROG
(lm x (lm y x))        ; => ⟨function⟩
((lm x x) (lm x x))  ; stx error! fn application undefined
```

A module declaring `#lang lam` may only write `lm` functions; using any other form results in an error. A Racket module may also serve as a *library*, as in the following LC module:

```
#lang racket                                     LC
(require lam) (provide (rename [lm λ][app #%app]))
(define-m (app e_fn e_arg)  (#%app e_fn e_arg))
```

LC imports `lm` from LAM and also defines `app`, which corresponds to single-argument function application. LC exports `lm` and `app` with new names, `λ` and `#%app`, respectively. The `#%app` in the output of `app` is core-Racket's function application form, though programmers need not write it explicitly. Instead, macro expansion implicitly inserts it before applied functions. This enables modifying the behavior of function application, as we do here by exporting `app` as `#%app`. Thus a program in the LC language looks like:

```
#lang lc                                    LC-PROG
((λ x (x x)) (λ x (x x))) ; => loop!
```

where λ corresponds to `lm` in LAM and applying a λ behaves according to `app` in LC. Running LC-PROG loops forever.

Figure 1 depicts compilation of a Racket program, which includes macro expansion. The Racket compiler first "reads" a program's surface text into a syntax object, which is a tree of symbols and literals along with context information, e.g.,
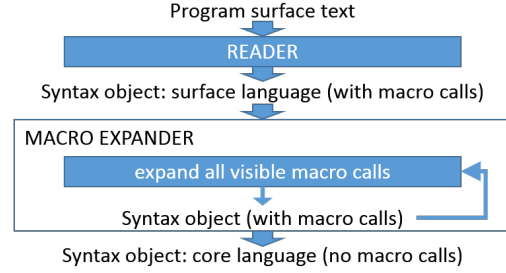

Figure 1. The Racket compiler's frontend

$$\tau ::= \tau \to \tau \quad e ::= x \mid \lambda x{:}\tau.e \mid e\,e \quad \Gamma ::= x{:}\tau, \ldots \quad \text{(types, terms)}$$

$$\frac{x{:}\tau \in \Gamma}{\Gamma \vdash x \,:\, \tau} \quad \text{(T-VAR)} \qquad \frac{\Gamma, x{:}\tau_1 \vdash e \,:\, \tau_2}{\Gamma \vdash \lambda x{:}\tau_1.e \,:\, \tau_1 \to \tau_2} \quad \text{(T-ABS)}$$

$$\frac{\Gamma \vdash e_1 \,:\, \tau_1 \to \tau_2 \qquad \Gamma \vdash e_2 \,:\, \tau_1}{\Gamma \vdash e_1\,e_2 \,:\, \tau_2} \quad \text{(T-APP)}$$

$$er(x) = x, \quad er(\lambda x{:}\tau.e) = \lambda x.er(e), \quad er(e\,e') = er(e)\,er(e') \quad \text{(erase)}$$

Implementation sketch:
```
#lang racket                                   STLC
(define-m (checked-λ ...)  ⟨when T-ABS⟩ er(———))
(define-m (checked-app ...)  ⟨when T-APP⟩ er(———))
```

Figure 2. Simply-typed λ-calculus

in-scope bindings and source locations. The macro expander then expands macro invocations in this syntax object according to macro definitions from the program's declared `#lang`. Macro expansion may reveal additional macro uses or even define new macros, so expansion repeats until no macro uses remain. Compilation terminates with a syntax error if expansion of any macro fails. The output of macro expansion contains only references to Racket's core syntax. This paper shows how to embed type checking within macro expansion.

3. A Typed λ-calculus Embedded Language

LC from section 2 implements the untyped λ-calculus. This section augments LC with types and type checking by transcribing formal type rules directly into its macro definitions, producing the simply-typed λ-calculus and demonstrating that Racket's macro infrastructure can be reused for type checking. Figure 2 presents standard STLC rules and a skeleton implementation. The macros also erase types in the surface language to accommodate the untyped Racket host.

3.1 Typed function application

Figure 3 presents `checked-app`, a macro that elaborates typed function application nodes into core Racket and also type checks the syntax tree ("v0" marks this initial version). Additional `#:with` and `#:when` conditions guard the macro's expansion. A pattern and expression follow a `#:with` and macro expansion continues only if the result of evaluating the latter produces a syntax object that matches the former. The first `#:with` uses a `compute-τ` function to compute the

```
1  (define-m (checked-app e_fn e_arg)  ; v0
2    #:with (→ τ_in τ_out) (compute-τ e_fn)
3    #:with τ_arg (compute-τ e_arg)
4    #:when (τ= τ_arg τ_in)
5    #:with ē_fn (erase-τ e_fn)
6    #:with ē_arg (erase-τ e_arg)
7    (add-τ (#%app ē_fn ē_arg) τ_out))
```

Figure 3. A type checking function application macro

```
1  (define (add-τ e τ)      (add-stx-prop e 'type τ))
2  (define (get-τ e)        (get-stx-prop e 'type))
3  (define (compute-τ e)    (get-τ (local-expand e)))
4  (define (erase-τ e)      (local-expand e))
5  (define (comp+erase-τ e)  ; get e's type, erase types
6    #:with ē (local-expand e) #:with τ (get-τ ē)
7    [ē τ])
8  (define (τ= τ_1 τ_2) (stx= τ_1 τ_2))
```

Figure 4. Helper functions for type checking

type of function e_{fn}, which must match pattern $(→ τ_{in} τ_{out})$.
The second `#:with` computes the type of argument e_{arg},
binding it to pattern variable $τ_{arg}$. Unlike the first `#:with`,
the $τ_{arg}$ pattern does not constrain the shape of e_{arg}'s type
but the following `#:when` asserts that $τ_{arg}$ and $τ_{in}$ satisfy
predicate $τ=$. The types in e_{fn} and e_{arg} are then erased (lines
5-6) before they are emitted in the macro's output (overlines
mark type-erased expressions, and core Racket forms). Fi-
nally, `add-τ` (line 7) "adds" $τ_{out}$ to the macro's syntax object
output. In summary, `checked-app` rewrites a typed function
application to an equivalent untyped one, along with its type.

3.2 Communicating macros

The organization of `checked-app` in figure 3 resembles a
combination of its T-APP and *erase* specification in figure 2.
Figure 4 defines some functions needed by `checked-app`,
which together establish a communication protocol between
type rule macros. The functions utilize *syntax properties*,
which are arbitrary key-value pairs stored with a syntax
object's metadata. For example, `checked-app` calls `add-τ`
to attach type information to its output, which in turn calls
`add-stx-prop` (figure 4, line 1) to associate a type $τ$ with
key `'type` on expression e. If all type rule macros follow
this protocol, then to compute an *arbitrary* expression's type,
we simply invoke that expression's macro and retrieve the
attached type from its output. In other words, *expanding an*
expression also type checks it.

We can call Racket's macro expander to invoke the de-
sired type checking macro but not in the standard manner.
Macro expansion typically rewrites *all* macro invocations in
a program at once (figure 1) and repeats this process until

```
1  (define-m (checked-app e_fn e_arg)  ; v1
2    #:with [ē_fn (→ τ_in τ_out)] (comp+erase-τ e_fn)
3    #:with [ē_arg τ_arg] (comp+erase-τ e_arg)
4    #:when (τ= τ_arg τ_in)
5    (add-τ (#%app ē_fn ē_arg) τ_out))
```

Figure 5. Revise fig 3 to compute and erase types together

there are no more macro calls. Such breadth-first expansion
is incompatible with type checking, however, which pro-
ceeds in a depth-first manner—a term is well-typed only
if its subterms are well-typed—but the `local-expand` [17]
function controls expansion in the desired way, expanding
just one syntax object without considering other parts of
the program. Thus `compute-τ` expands its argument with
`local-expand` (figure 4, line 3) and then retrieves its type.

The `checked-app` macro uses `erase-τ` to produce syntax
without type annotations. If all type rule macros follow this
protocol then *expanding an expression also erases its types*.
Separate calls to `compute-τ` and `erase-τ`, however, unnec-
essarily expands syntax twice. The `comp+erase-τ` function
(lines 5-7) eliminates this redundancy and figure 5's revised
`checked-app` uses this function. In general, we carefully
avoid extraneous expansions while type checking so as not
to change the algorithmic complexity of macro expansion.

Finally, type checking requires a notion of type equality.
We cannot compute mere symbolic equality since types are
renamable linguistic constructs:

```
(require (rename [→ a])) (τ= (a s t) (→ s t)) ; ⇒ true
```

If we represent types with syntax objects, however, *type*
equality is syntax equality and we can reuse Racket's knowl-
edge of the program's binding structure (`stx=` in figure 4 line
8) to compute type equality in a straightforward manner.

3.3 Type environments and type checking $λ$

Figure 6 implements the $→$ type (lines 1-2) as a macro that
matches on an input and output type and expands to an appli-
cation of an internal function that errors at runtime (there are
no base types for now, see §4.3). The `checked-λ` macro re-
quires a type annotation on its parameter (line 4), separated
with `:`. This macro resembles `checked-app`, except a new
`comp+erase-τ/ctx` function replaces `comp+erase-τ`. Since
the $λ$ body may reference x, `comp+erase-τ/ctx` computes
the body's type in a type context containing x and its type,
given as the second argument.

So far, `checked-app` and `checked-λ` correspond to T-APP
and T-ABS from figure 2, respectively. To implement T-
VAR, i.e., type environments, `comp+erase-τ/ctx` defines a
local macro with `let-macro` (figure 6, line 11) and expands
an expression e in the scope of this new macro. The local
macro is named x and expands to a fresh y that has the de-
sired type $τ$ attached (observe the nested gray highlights).

```
1  (define →intrnl (λ _ (ERR ''no runtime types'')))
2  (define-m (→ τin τout) (→intrnl τin τout))
3
4  (define-m (checked-λ [x : τin] e)  ; v0
5    #:with [x̄ ē τout] (comp+erase-τ/ctx e [x τin])
6    (add-τ (λ̄ (x̄) ē) (→ τin τout)))
7
8  (define (comp+erase-τ/ctx e [x τ])
9    #:with (λ̄ (x̄) ē)
10           (local-expand ; y fresh
11             (λ̄ (y) (let-macro [x (add-τ y τ)] e)))
12    #:with τout (get-τ ē)
13    [x̄ ē τout])
```

Figure 6. Type checking λ and → macros

```
1  (define-m (→ τin ... τout) (→intrnl τin ... τout))
2
3  (define-m (checked-app efn earg ...)  ; v2
4    #:with [ēfn (→ τin ... τout)] (comp+erase-τ efn)
5    #:with ([ēarg τarg] ...) (comp+erase-τ⃗ (earg ...))
6    #:fail-unless (τ⃗= (τarg ...) (τin ...))
7       (fmt ''~a: Fn args have wrong types: expected: ~a, got: ~a''
8         (src this-stx) (τin ...) (τarg ...))
9    (add-τ (#%app ēfn ēarg ...) τout))
10
11 (define-m (checked-λ ([x : τin] ...) e)  ; v1
12    #:with [xs̄ ē τout]
13           (comp+erase-τ/ctx e ([x τin] ...))
14    (add-τ (λ̄ xs̄ ē) (→ τin ... τout)))
```

Figure 7. Multi-arity functions and error checking

As a result, while expanding e, a reference to x (with no type information) becomes a reference to y (with type information). To avoid unbound y errors during expansion, a (Racket) λ wraps the let-macro before expansion. Finally, comp+erase-τ/ctx returns a tuple of post-expansion y (as x̄), the type-erased ē, and its type τout. Effectively, defining a local macro inserts a binding indirection level during macro expansion, enabling the insertion of the desired type information on variable references. Thus T-Var is implemented, reusing the compile-time macro environment as the type environment. This completes our simply-typed language.

3.4 A few practical matters

We have implemented a basic λ-calculus; however, we wish to implement practical languages. This subsection shows how to extend our language with some practical features.

Multiple arguments Figure 7 revises our language to support multiple arguments and useful error messages. An ellipsis pattern (...) matches zero-or-more of the preceding element. If that preceding element binds pattern variables, ellipses must follow later references to those variables, e.g., the revised → macro (line 1) matches zero-or-more input arguments τin and ellipses follow τin in its output. The other forms are extended similarly. The checked-λ macro uses a slightly modified comp+erase-τ/ctx (line 13) that accepts multi-element contexts. In checked-app (line 5), the "vector" f⃗ notation denotes f mapped over its input list.

Error messages The checked-app in figure 5 reports type errors as syntax errors but a better message should indicate the error's location and the computed and expected types. The checked-app in figure 7 reports such a message using a #:fail-unless condition (lines 6-8) to produce a message from a printf-style string and arguments (this-stx is the current input syntax, analogous to the OO "this"). All our languages strive to report accurate messages in the manner of figure 7, though the paper may not always show this code.
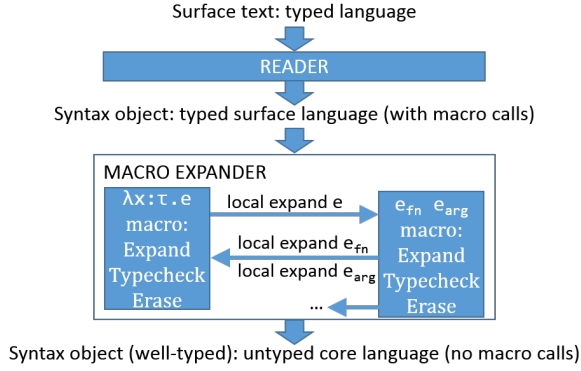
```
1  (define-m #%type (#%typeintrnl))
2  (define (valid-τ? τ) (τ= (compute-τ τ) #%type))
3
4  (define-m (checked→ τ ...)
5    #:fail-if (nil? (τ ...))(ERR ''→ requires >=1 args'')
6    #:fail-unless (valid-τ?⃗ (τ ...))
7          (fmt ''→ given invalid types: ~a'' (τ ...))
8    (add-τ (→intrnl τ ...) #%type))
9
10 (define-m (checked-λ ([x : τin] ...) e)  ; v2
11    #:fail-unless (valid-τ?⃗ (τin ...))
12          (fmt ''λ given invalid types: ~a'' (τin ...))
13    #:with [xs̄ ē τout]
14           (comp+erase-τ/ctx e ([x τin] ...))
15    (add-τ (λ xs̄ ē) (→ τin ... τout)))
```

Figure 8. Checking type well-formedness

Type well-formedness Our language so far checks the types of terms but does not check whether programmer-written types are valid, e.g., one may write (λ ([x : (→)]) x) or (λ ([x : Undef]) x). Applying these functions result in type errors but the invalid types should be reported before then. Some type checkers validate types during parsing; this is undesirable for our purposes since it prevents defining types not expressible with a grammar. Instead, we use kinds.

To check kinds, we use *the same type checking technique* from our term-checking macros. Figure 8 defines a single kind named #%type and all types are tagged with this kind (e.g., line 8). Thus, → and λ may validate their input types with valid-τ? (lines 6-7, 11-12). The use of the macro expander to validate types also differentiates when a type is

Figure 9. Macro-based typed language implementations

Figure 10. Interleaved typechecking and erasure rules

undefined, rather than malformed. Ultimately, the previous examples now produce type errors:

```
(λ ([x : (→)]) x)    ; TYERR: → requires >= 1 args
(λ ([x : Undef]) x)  ; TYERR: unbound id Undef
```

4. A Metalanguage For Typed Languages

4.1 Interleaved type checking and rewriting

Section 3's STLC implementation reveals a synergy between macro expansion and type checking in that Racket's macro infrastructure can be reused to also check and erase types during its program traversal. Figure 9 refines figure 1 to incorporate this reuse. This organization further suggests a reformulation of figure 2's rules to combine typechecking and erasure, shown in figure 10. A new $\Gamma \vdash e \gg \bar{e} : \tau$ rule reads "in context Γ, e erases to $\bar{e}$ and has type τ", where contexts consist of variable "erasures", e.g., TE-ABS inserts a binding indirection level in the context in order to add type information for variables and checks a λ body in this context. These rules straightforwardly correspond to our macro-based type system implementation in section 3, where $\Gamma \vdash e \gg \bar{e} : \tau$ is *implemented* as "in context Γ, e *expands* to $\bar{e}$, with type τ attached". As this paper focuses on implementation, we do not formally study these new typing rules, though they do suggest how to further improve our approach to implementing typed embedded languages.

Figure 11. The STLC implemented with TURNSTILE

4.2 The TURNSTILE metalanguage

Section 3 demonstrates that a macro system's infrastructure can be reused to implement typechecking. Deploying such an approach, however, requires writing macro-level code to embed type rules into macro definitions despite the resemblance of this code to its mathematical specification. This section introduces TURNSTILE, a Racket DSL for creating practical embedded languages that abstracts the macro-level ideas and insights from the previous section into linguistic constructs at the level of types and type systems.

Specifically, TURNSTILE enables writing rules using the syntax from figure 10 but with bidirectional [38] "synthesize" ($\Rightarrow$) and "check" ($\Leftarrow$) arrows replacing the colon, to further clarify inputs and outputs. Figure 11 reimplements STLC with TURNSTILE. TURNSTILE repackages all the infrastructure from section 3 as convenient abstractions, e.g., `define-type-constructor` (d-t-c) on line 2 and the subsequent `define-typerule`s (d-t) implementing `#%app` and λ.

TURNSTILE's syntax further demonstrates the connection between specification and implementation enabled by our macro-based approach. Though programmers may now write in a declarative syntax, STLC's implementation has not changed as TURNSTILE's abstractions are mere syntactic sugar for the macros from section 3. For example, $\Rightarrow$ abbreviates `#:with` used with `comp+erase-τ` and thus figure 11, line 4 exactly corresponds to figure 7, line 4. Similarly, $\Leftarrow$ abbreviates `#:with`, `#:fail-unless`, and `τ=` so figure 11, line 5 corresponds to figure 7, lines 5-8. Finally, $\Rightarrow$ below the conclusion line corresponds to `add-τ` as in figure 7, line 9 (crossing the conclusion line inverts the yellow and gray positions of $\Rightarrow$). The λ d-t's premise computes e's type in a type context containing the variables to the left of $\vdash$ (figure 11, line 9). The λ input pattern (line 8) utilizes annotations asserting that x is an identifier and τ_{in} is a valid type.

In general, a d-t resembles a figure 10 rule except the conclusion is split into its inputs and outputs—the (yellow) pattern(s) (and $\gg$) that begin a definition, and (gray) syntax following the conclusion line, respectively—such that the definition (and variable scoping) reads top-to-bottom. Fig-

ures 10 and 11 additionally differ because d-ts do not explicitly thread through a "Γ", a consequence of reusing Racket's scoping for the type environment. Thus Turnstile programmers only write *new* type environment bindings in d-ts, analogous to let; existing bindings are implicitly available according to standard lexical scope behavior.

Viewed as type rules, figure 11 appears to be missing the ⇐ rules. While a programmer may write explicit ⇐ rules (see §6), in their absence, TURNSTILE uses this default:

```
(define-typerule e ⇐ τ ≫
   [⊢ e ≫ e̅ ⇒ τₑ]
   [τₑ = τ]
   -----------------------
   [⊢ e̅])
```

This implicit definition corresponds to figure 7, lines 5-8. The first and last lines again comprise the input and output components of the rule's "conclusion", respectively, with the "expected" type now a part of the input pattern matching.

Though TURNSTILE programmers may implement type rules in a declarative style, such a style is likely insufficient for creating practical languages. Therefore, all the macro features from section 3 are also available to a d-t definition, giving TURNSTILE type rules access to the full power of Racket's macro system. For example, a programmer may add #:fail-unless error messages as in figure 7. Here is a refined #%app that further differentiates arity errors:

```
(define-typerule (#%app e_fn e_arg ...) ≫
   [⊢ e_fn ≫ e̅_fn ⇒ (→ τ_in ... τ_out)]
   #:fail-unless (len= (e_arg ...) (τ_in ...))
      (fmt ''~a: wrong number of args: expected ~a, got ~a''
         (src this-stx) (len (τ_in ...)) (len (e_arg ...)))
   [⊢ e_arg ≫ e̅_arg ⇐ τ_in] ...
   ----------------------------------
   [⊢ (#%app e̅_fn e̅_arg ...) ⇒ τ_out])
```

4.3 Reusing a Type System

TURNSTILE type rules from one language may be reused in the implementation of another. Though STLC implements function application and λ, it defines no base types and thus no well-typed programs. We next add integers and addition but instead of revising STLC, we reuse its rules in a new language, analogous to section 2. Specifically, STLC+PRIM in figure 12 uses STLC as a library, importing and re-exporting its type rules with extends. To STLC's definitions, STLC+PRIM adds an Int base type (line 2), a + primop (line 3), and integer literals (lines 4-7). Just as the macro expander inserts #%app before applied functions, it also wraps literals with #%datum, whose behavior is overridden in figure 12 to add types to integers. With STLC+PRIM, we can now write well-typed programs.

```
1  #lang turnstile    (extends stlc)      STLC+PRIM
2  (define-base-type Int)
3  (define-primop + : (→ Int Int Int))
4  (define-typerule (#%datum n) ≫
5     #:fail-unless (int? n) (ERR ''Unsupported datum'')
6     -----------------------------------
7     [⊢ (#%datum n) ⇒ Int])
```

Figure 12. A language extending STLC with integers

5. A Series of Core Languages

To confirm that our approach to typed languages handles a variety of type systems, we implemented a series of textbook core languages [37]. This section describes a few examples.

5.1 Types that bind: existential types

Figure 13 depicts EXIST, a language with existential types; it reuses records and variants from another language. The #:bvs option (line 2) specifies that an ∃ type binds one variable and thus has surface syntax (∃ (X) τ_body).

Figure 4 (line 8) introduced type equality as structural equality of syntax objects. Type equality of quantified types, however, should include alpha equivalence. While other systems commonly convert to alternate representations such as de Bruijn indices [6] to implement this behavior, our use of syntax objects for types remains sufficient since these objects contain knowledge of the program's binding structure. Thus the τ= used by TURNSTILE looks like:

```
(define [(τ= (C₁ Xs τ₃) (C₂ Ys τ₄))
         (and (τ= C₁ C₂) (τ= (subst Ys Xs τ₃) τ₄))]
        [else ————]) ; structural traversal
```

This updated τ= specifies multiple input patterns. The first clause matches binding types where equality of such types with the same constructor is equivalent to renaming parameter references to the same name and recursively comparing the resulting body for equality. Otherwise, types are structurally compared. A subst function performs this renaming:

```
(define (subst v x e)
   (if (and (id? e) (binds? x e)) v ——— ; else traverse e))
```

Specifically, (subst v x e) replaces occurrences of x in e with v, where binds? determines "occurrence" by examining lexical information in the syntax objects. Thus substitution is a structural traversal and no renaming is necessary.

The pack and open macros use τ= and subst: pack assigns a term e an existential type (∃ (X) τ_body), where e has concrete type equal to replacing X in τ_body with τ_hide; dually, open binds x to an existentially-typed e_packed's value, type variable X to e_packed's hidden type, and then checks an expression e in the context of X and x. To the left of ⊢ (figure 13, line 9) is two environments: a list of type variables and the standard environment for term variables. The (∃ (Y) τ_bod) type of e_packed is "opened", so x has type τ_bod but with occurrences of the existentially-bound Y (not in scope in e)

```
1  #lang turnstile   (extends stlc+reco+var)    EXIST
2  (define-type-constructor ∃ #:bvs = 1)
3  (define-typerule (pack [τ_hide^type e] as (∃(X) τ_body)) ≫
4    [⊢ e ≫ ē ⇐ (subst τ_hide X τ_body)]
5    ---------------------------------
6    [⊢ ē ⇒ (∃ (X) τ_body)])
7  (define-typerule (open [x^id e_packed] with X^id in e) ≫
8    [⊢ e_packed ≫ ē_packed ⇒ (∃ (Y) τ_bod)]
9    [(X)([x ≫ x̄ : (subst X Y τ_bod)]) ⊢ e≫ ē ⇒τ]
10   ---------------------------------
11   [⊢ (let ([x̄ ē_packed]) ē) ⇒ τ])
```

Figure 13. A language with existential types

replaced with its "opened" X name. Here is a typical counter example (×, rcrd, and prj correspond to records):

```
#lang exist
(define COUNTER
 (pack [Nat (rcrd [new = 1][inc = add1]
                  [get = (λ ([x : Nat]) x)])] as
  (∃ C (× [new : C][inc :(→ C C)][get :(→ C Nat)])))))
(open [c COUNTER] with Count in
 (+ ((prj c get) ((prj c inc) (prj c new))) ; => 2
    (add1 (prj c new)))); TYERR: expected type Nat, got Count
```

5.2 Subtyping and enhanced modularity

Figure 14 presents STLC+SUB, a language with subtyping that reuses parts of STLC+PRIM from figure 12 but adds new base types and redefines #%datum and + with these types. One might not expect STLC+SUB to reuse type rules that do not consider subtyping. However, TURNSTILE exposes hooks for common type operations and implements type checking in terms of these hooks, enabling better reuse. For example, τ= in figure 5, line 4 is actually an overridable "type check relation" (initially set to τ=). These language-level hooks are implemented with Racket *parameters* [20]. Thus STLC+SUB defines a new τ<: predicate and installs it as the $\boxed{\tau\sqsubseteq}$ type check relation (we oval-box parameter names), enabling reuse of #%app and λ from STLC.

TURNSTILE pre-defines parameters like $\boxed{\tau\sqsubseteq}$ and $\boxed{\tau\ eval}$; the latter is called before attaching types to syntax. Each language may also define new parameters, e.g., STLC+SUB additionally defines $\boxed{join}$ and uses it in conditionals.

5.3 Defining types and kinds

Our implementations macro-expand a term to type check and erase its types. We can check kinds the same way: expanding a type kind checks and erases kinds. The kind erasing may cause problems, however, since a type judgement may use both types and kinds. Nevertheless, TURNSTILE can define a kind system like in F_ω. To address the problem, figure 15 reformulates some F_ω rules with our $\gg$ relation. Specifically, T-TABS and K-ALL erase a $\forall$'s kind annotation, but "save" it with $\star$, now a kind constructor, in the same manner that $\rightarrow$

```
1  #lang turnstile                                STLC+SUB
2  (extends stlc+prim #:except #%datum +)
3  (define-base-types Top Num Nat)
4  (define-typerule #%datum ————————)
5  (define-primop + : (→ Num Num Num))
6  (define-primop add1 : (→ Int Int))
7  (define (τ<: τ_1 τ_2)
8    (or (τ= τ_1 τ_2)
9      (syntax-parse (τ_1 τ_2)
10     [(_ Top) true]
11     [(_ Num) (τ<: τ_1 Int)]
12     [(_ Int) (τ<: τ_1 Nat)]
13     [((→ τ_i1 ... τ_o1) (→ τ_i2 ... τ_o2))
14       (and (→⃗<: (τ_i2 ...) (τ_i1 ...)) (τ<: τ_o1 τ_o2))]
15     [else false])))
16 (set-�τ⊑ τ<:)  ; no need to redefine #%app or other rules
```

Figure 14. A simply-typed language with subtyping

$$e ::= x \mid \lambda x{:}\tau.e \mid e\,e \mid \Lambda X{::}\kappa.e \mid e\,\tau \quad \text{(terms with types)}$$
$$\tau ::= X \mid \tau\rightarrow\tau \mid \forall X{::}\kappa.\tau \mid \lambda X{::}\kappa.\tau \mid \tau\,\tau \quad \text{(types with kinds)}$$
$$\bar{e} ::= \bar{x} \mid \lambda\bar{x}.\bar{e} \mid \bar{e}\,\bar{e}, \quad \kappa ::= \star\kappa\ldots \mid \kappa \Rightarrow \kappa \quad \text{(typeless terms, kinds)}$$
$$\bar{\tau} ::= \overline{X} \mid \bar{\tau}\rightarrow\bar{\tau} \mid \forall\overline{X}.\bar{\tau} \mid \lambda\overline{X}.\bar{\tau} \mid \bar{\tau}\,\bar{\tau} \quad \text{(kindless types)}$$
$$\Gamma ::= b,\ldots \quad b ::= X \gg \overline{X}{::}\kappa \mid x \gg \bar{x}{:}\tau{::}\kappa \quad \text{(contexts)}$$

$$\frac{\Gamma, X \gg \overline{X}{::}\kappa \vdash e \gg \bar{e} : \bar{\tau} :: \kappa}{\Gamma \vdash \Lambda X{::}\kappa.e \gg \bar{e} : \forall\overline{X}.\bar{\tau} :: \star\kappa} \quad \text{(T-TABS)} \quad \boxed{Typing}$$

$$\frac{\Gamma \vdash e \gg \bar{e}{:}\forall\overline{X}.\bar{\tau}_2{::}\star\kappa \quad \Gamma \vdash \tau \gg \bar{\tau}{::}\kappa \quad \Gamma \vdash \bar{\tau}_2{::}\kappa_2}{\Gamma \vdash e\,\tau \gg \bar{e} : \bar{\tau}_2[\overline{X} \leftarrow \bar{\tau}] :: \kappa_2} \quad \text{(T-TAPP)}$$

$$\frac{X \gg \overline{X}{::}\kappa \in \Gamma}{\Gamma \vdash X \gg \overline{X} :: \kappa} \quad \text{(K-VAR)} \quad \boxed{Kinding}$$

$$\frac{\Gamma, X \gg \overline{X}{::}\kappa \vdash \tau \gg \bar{\tau} :: \star_{-}\ldots}{\Gamma \vdash \forall X{::}\kappa.\tau \gg \forall\overline{X}.\bar{\tau} :: \star\kappa} \quad \text{(K-ALL)}$$

Figure 15. Some F_ω rules using figure 10's $\gg$ relation

"saves" a λ's type annotations. T-TAPP then checks that its argument type has a kind matching the saved annotation.

Figure 16 implements FOMEGA utilizing figure 15's insights: it introduces a "kind" category of syntax, defines $\Rightarrow$ and $\star$ kinds, and directs the $\forall$ type to construct its kind with a $\star$ "arrow". Lines 8-9 connect "kinds" and "types" where line 8 enables reuse of previously-defined types, and line 9 redefines "well-formed" types. Finally, the Λ rule type checks its body in its type variable's context, and the inst rule instantiates an expression e at type τ by computing e's $\forall$ type and that type's kind ($\star\kappa$), and checking that τ has kind κ.

5.4 Reusing languages

Table 1 summarizes extensions and reuse in fourteen core language implementations. A row and color represents a language and features are in columns. A diamond marks

```
1  #lang turnstile                              FOMEGA
2  (extends stlc+prim)  (reuse τ= #:from exist)
3  (define-stx-category kind)
4  (define-kind-constructor ⇒ #:arity >= 1)
5  (define-kind-constructor ⋆ #:arity >= 0)
6  (define-type-constructor ∀ #:bvs = 1 #:arrow ⋆)
7  ; link types and kinds
8  (set-kind? (λ (k) (or (#%type? k) (kind? k))))
9  (set-type? (λ (t) (and (kind? t) (not (⇒? t)))))
10 ; ...
11 (define-typerule (Λ [X^id :: κ^kind] e) ≫
12   [([X ≫ X̄ :: κ]) () ⊢ e ≫ ē ⇒ τ̄_e]
13   -----------------------------------
14   [⊢ ē ⇒ (∀ ([X̄ : κ]) τ̄_e)])
15 (define-typerule (inst e τ) ≫
16   [⊢ e ≫ ē ⇒ (∀ X̄ τ̄_body) (⇒ (⋆ κ))]
17   [⊢ τ ≫ τ̄ ⇐ κ]
18   -----------------------------------
19   [⊢ ē ⇒ (subst τ̄ X̄ τ̄_body)])
```

Figure 16. A language with higher-order polymorphism

a feature's first implementation and down-column appearances of the feature's color indicates reuse. Thus single-color columns and multi-color rows indicate *abundant* reuse. For example, all languages share the same λ; also, languages with basic types share a τ= while those with binding types use an extended version. A $\oplus$ marks feature extension (a dotted line connects non-adjacent-row extensions). For example, "subtyping" extends #%datum from "stlc+prim"; also, $F_{<:}$ extends and combines subtyping from one language and $\forall$ from system F to implement bounded polymorphism.

Table 2 summarizes implementation sizes from table 1. Each column represents a different implementation of the same language: the first uses TURNSTILE; the second uses TURNSTILE but does not import other implementations; and the third uses plain Racket. Though the last two columns are estimates (2 sig. figs.)—we did not implement every permutation of every language—they still indicate the degree of reuse. Roughly, the second and first column difference represents the degree to which type rules are reused across many languages, analogous to single-color columns in table 1. Such reuse would be difficult to achieve with conventional type checker implementations. The third and first column difference indicates the degree to which TURNSTILE captures common patterns used to implement type checkers.

5.5   More than types: a type-and-effect system

Table 1's languages mostly use a typical $\Gamma \vdash e : \tau$ relation though TURNSTILE is not limited to this relation. Rather, programmers may specify propagation of any number of arbitrary properties. For example, figure 17 presents EFFECT, a language with a basic type and effect system [34]. The language adds Void and Ref types, and ref, deref, and :=

```
1  #lang turnstile                              EFFECT
2  (extends stlc+prim #:except #%app λ)
3  (define-base-type Void)
4  (define-type-constructor Ref #:arity = 1)
5  (define-typerule (ref e) ≫
6    [⊢ e ≫ ē (⇒ : τ) (⇒ :ν π)]
7    -----------------------------------
8    [⊢ (box ē) (⇒ : (Ref τ))
9              (⇒ :ν (∪ (src (ref e)) π))])
10 (define-typerule (deref e) ≫
11   [⊢ e ≫ ē (⇒ : (Ref τ)) (⇒ :ν π)]
12   -----------------------------------
13   [⊢ (unbox ē) (⇒ : τ) (⇒ :ν π)])
14 (define-typerule (:= e_ref e) ≫
15   [⊢ e_ref ≫ ē_ref (⇒ : (Ref τ_ref)) (⇒ :ν π_ref)]
16   [⊢ e ≫ ē (⇐ : τ_ref) (⇒ :ν π)]
17   -----------------------------------
18   [⊢ (set-box! ē_ref ē) (⇒ : Void) (⇒ :ν (∪ π_ref π))])
19 (define-typerule (#%app e_fn e_arg) ≫
20   [⊢ e_fn ≫ ē_fn (⇒ : (→ τ_in τ_out) (⇒ :ν π_app))
21              (⇒ :ν π_fn)]
22   [⊢ e_arg ≫ ē_arg (⇐ : τ_arg) (⇒ :ν π_arg)]
23   -----------------------------------
24   [⊢ (#%app ē_fn ē_arg) (⇒ : τ_out)
25              (⇒ :ν (∪ π_fn π_arg π_app))])
26 (define-typerule (λ [x^id : τ_in^type] e) ≫
27   [[x ≫ x̄ : τ_in] ⊢ e ≫ ē (⇒ : τ_out) (⇒ :ν π)]
28   -----------------------------------
29   [⊢ (λ̄ (x̄) ē) (⇒ : (→ τ_in τ_out) (⇒ :ν π))])
```

Figure 17. A basic side effect analysis.

type rules for allocation of, dereference of, and assignment to reference cells, respectively (box is Racket's ref cells). In addition to types, the language tracks source locations (π) of ref allocations (line 9). The ref rule exhibits new syntax: instead of a type to the right of $\Rightarrow$, a programmer may write multiple $\Rightarrow$ arrows matching multiple properties. Thus ref specifies that expansion of e (line 6) computes both a type (keyed on :) and a set of locations π (keyed on :$_\nu$). The key symbols match the user-specified symbols below the conclusion line. The := rule uses both $\Leftarrow$ (for the type) and $\Rightarrow$ (for the locations) simultaneously (line 16).

EFFECT contrasts with table 1's languages in that it cannot reuse #%app and λ due to its incompatible type relation. (It does reuse some types and type operations.) The new #%app and λ rules show that both terms and types carry the :$_\nu$ property. Specifically, λ propagates :$_\nu$ to function types (line 29), expressed with a *nested* $\Rightarrow$ (like the double-$\Rightarrow$ syntax for kinds from figure 16), because evaluating a λ does not trigger allocations in its body. Applying a function does evaluate the body, so #%app transfers locations from the function type (line 20) to the application term (line 25).

Table 1. Implemented Languages

feature / lang name	λ	app	→	τ=	τeval	Int	+	datum	Bool if letrec begin define alias	tup type	tup	recrd	variant	List	Ref	Top Nat Num	τ<:	join	μ	∃	∀	Λ inst	tyapp tyλ	invoke	kind
stlc	◇	◇	◇	◇	◇																				
stlc + prim						◇	◇	◇																	
extended stlc								⊕	◇																
tuples										◇	◇														
records + variants										⊕		◇	◇												
lists														◇											
reference cells															◇										
subtyping							◇	⊕								◇	◇								
subtyping + records								⊕									⊕	◇							
(iso) recursive				⊕															◇						
existential																				◇					
system F																					◇	◇			
F<:																	⊕				⊕	◇		◇	
F omega					⊕																⊕	◇	◇		◇

◇ = implemented here   ⊕ = extends the above

A unique color represents each language. The features in each language (row) are colored according to the language where they are defined.

Language Name	w/ TURNSTILE	no reuse*	w/ Racket*
stlc	32	32	
stlc+prim	23	55	
extended stlc	143	200	
tuples	32	230	
records+variants	171	400	
lists	73	470	~1100
reference cells	29	500	to
subtyping	107	160	~1300
sub+records	50	610	LoC
(iso) recursive	27	260	
existential	69	470	
system F	28	83	
$F_{<:}$	89	700	
F_ω	112	190	
* = estimate			

Table 2. Line count comparisons of table 1's languages

6. A Full-sized Language

To show that TURNSTILE scales to real-world type systems, we created MLISH, an ML-like language with local type inference, recursive user-defined algebraic data types, pattern matching, and Haskell-style type classes [46], along with "batteries included" features like efficient data structures, mutable state, generic sequence comprehensions, I/O, and concurrency primitives. MLISH also demonstrates how TURNSTILE easily incorporates type-system-directed program transformations. This section explains a few features.

Local type inference MLISH aims to follow Pierce and Turner's empirical inference guidelines [38]. Specifically, programmers need not write most annotations and instantiations except top-level function signatures, which are useful as documentation, and some λ annotations, which are rare.

Figure 18 sketches basic type inference in λ and #%app. Multiple clauses comprise λ, whose input patterns are checked in order. The first clause matches unannotated λs whose context determines its type, indicated with $\Leftarrow$ (line 5). The second matches annotated λs with implicitly bound type variables, computes these variables, and then recursively invokes the λ rule (indicated with $\succ$) with explicit type variables. In this manner, a surface language with implicit type variables rewrites to one with explicit binders, reusing the macro system for the type-system-directed rewrite. Finally, the third clause matches λs with explicit type variable binders; it resembles λ from figure 11. An MLISH define for top-level functions uses λ, splitting a definition into a runtime component and a macro that adds type information:

```
(define-typerule (define (f [x : τ] ... → τ_out) e) ≫
  #:with (X ...) (free-tyvars (τ ...))
  #:with τ_f (∀ (X ...) (→ τ ... τ_out))
  [⊢ (λ (x ...) e) ≫ ē_λ ⇐ τ_f]
  --------------------------------
  [≻ (define f_intrnl ē_λ)
     (define-m f (add-τ f_intrnl τ_f))])
```

To implement a $\Leftarrow$ type rule, e.g., figure 18 lines 5-8, MLISH propagates "expected type" information from an ex-

```
1  #lang turnstile                          MLISH
2  ; ....
3  (define-typerule λ
4    ; no annotations, use expected type
5    [(λ (x^id ...) e) ⇐ (∀ (X ...) (→ τ_in ... τ_out)) ≫
6     [(X ...) ([x ≫ x̄ : τ_in] ...) ⊢ e ≫ ē ⇐ τ_out]
7     ------------------------------------
8     [⊢ (λ̄ (x̄ ...) ē)]]
9    ; variable annotations, with free tyvars
10   [(λ ([x^id : τ_in] ...) e) ≫
11    #:with (X ...) (free-tyvars (τ_in ...))
12    ------------------------------------
13    [≻ (λ {X ...} ([x : τ_in] ...) e)]]
14   ; variable annotations, explicit tyvar binders
15   [(λ {X ...} ([x^id : τ_in] ...) e) ≫
16    [(X ...) ([x ≫ x̄ : τ_in] ...) ⊢ e ≫ ē ⇒ τ_out]
17    ------------------------------------
18    [⊢ (λ̄ (x̄ ...) ē) ⇒ (∀ (X ...) (→ τ_in ... τ_out))]])
19  (define-typerule #%app
20    ; infer polymorphic instantiation, with expected type
21    [(#%app e_fn e_arg ...) ≫
22     #:when (has-expected-τ this-stx)
23     #:with τ_expct (get-expected-τ this-stx)
24     [⊢ e_fn ≫ ē_fn ⇒ (∀ Xs (→ τ_X ...))]
25     [⊢ e_arg ≫ ē_arg ⇒ τ_arg] ...
26     #:with (τ ...) (solve Xs (τ_arg ... τ_expct) (τ_X ...))
27     ------------------------------------
28     [≻ (#%app {τ ...} ē_fn ē_arg ...)]]
29    ; infer polymorphic instantiation, no expected type
30    [(#%app e_fn e_arg ...) ≫
31     [⊢ e_fn ≫ ē_fn ⇒ (∀ Xs (→ τ_inX ... τ_outX))]
32     [⊢ e_arg ≫ ē_arg ⇒ τ_arg] ...
33     #:with (τ ...) (solve Xs (τ_arg ...) (τ_inX ...))
34     ------------------------------------
35     [≻ (#%app {τ ...} ē_fn ē_arg ...)]]
36    ; explicit instantiation of polymorphic function
37    [(#%app {τ^type ...} e_fn e_arg ...) ≫
38     [⊢ e_fn ≫ ē_fn ⇒ (∀ Xs (→ τ_X ...))]
39     #:with (τ_in ... τ_out) (subst⃗ (τ ...) Xs (τ_X ...))
40     [⊢ e_arg ≫ ē_arg ⇐ τ_in] ...
41     ------------------------------------
42     [⊢ (#%app ē_fn ē_arg ...) ⇒ τ_out]])
```

Figure 18. Type inference in MLISH #%app and λ

pression's context by attaching a syntax property *before* expansion, making the information available while typechecking that expression. A $\Leftarrow$ type rule's input matches on this expected type (line 5), and also implicitly attaches it to the output syntax (line 8). A non-$\Leftarrow$ type rule may also inspect the expected type, as with #%app. Specifically, the first #%app clause extracts the expected type (lines 22-23) and uses it to solve for the type variables (line 26). The clause then re-

```
1  (define-m (define-type (Ty X ...)
2                         (Constr [fld : τ_X] ...) ...)
3    (define-type-constructor Ty
4      #:arity = (len (X...))
5      #:extra ((Constr [fld τ_X] ...) ...))
6    (struct Constr_intrnl (fld ...)) ...
7    (define-typerule (Constr e_arg ...) ≫
8      #:with C (add-τ Constr_intrnl
9                  (∀ (X ...) (→ τ_X ... (Ty X ...))))
10     ------------------------------------
11     [≻ (#%app C e_arg ...)]) ... )
12 (define-typerule (match e_m with [C x ... -> e]...+) ≫
13    [⊢ e_m ≫ ē_m ⇒ τ_m]
14    #:with [(Constr [fld τ_fld]...) ...] (get-extra τ_m)
15    #:fail-unless (set= (C ...) (Constr ...))
16                   (fmt ``not enough clauses, missing ~a''
17                        (set-diff (C ...) (Constr ...)))
18    [[x ≫ x̄ : τ_fld] ... ⊢ e ≫ ē ⇒ τ] ...
19    #:when (same-τ? (τ ...))
20    ------------------------------------
21    [⊢ (let ([v̄ ē])
22        (cond [(Constr? v̄)
23               (let ([x̄ (get-fld v̄)] ...) ē)] ...))
24       ⇒ (first (τ ...))])
```

Figure 19. Defining types and pattern matching in MLISH

```
1  (define-m (define-tc (Cls X ...) op_generic : τ_op)
2    (define-m (Cls X ...) [op_generic : τ_op])
3    (define-typerule (op_generic e ...) ⇐ τ_o ≫
4       [⊢ e ≫ ē ⇒ τ] ...
5       #:with op_concrete (lookup op_generic (→ τ... τ_o))
6       ------------------------------------
7       [≻ (#%app op_concrete ē ...)]) )
8  (define-m (define-instance (Cls τ...) op_gen op_c)
9    #:with [op : τ_concrete] (local-expand (Cls τ ...))
10   #:when (equal? op_gen op)
11   [⊢ op_c ≫ op_concrete ⇐ τ_concrete]
12   #:with op_mang (mangle op τ_concrete)
13   ------------------------------------
14   [≻ (define-m op_mang (add-τ op_concrete τ_concrete))])
15 (define-typerule #%app
16   ; ....
17   [(#%app {τ^type ...} e_fn e_arg ...) ≫
18    [⊢ e_fn ≫ ē_fn ⇒ (∀ Xs (=> TC (→ τ_X ...)))]
19    #:with (τ_in ... τ_out) (subst (τ ...) Xs (τ_X ...))
20    [⊢ e_arg ≫ ē_arg ⇐ τ_in] ...
21    #:with [op_generic : τ_generic] TC
22    #:with τ_concrete (subst (τ ...) Xs τ_generic)
23    #:with op_concrete (lookup op_generic τ_concrete)
24    ------------------------------------
25    [⊢ (#%app ē_fn op_concrete ē_arg ...) ⇒ τ_out]])
```

Figure 20. Type classes in MLISH

cursively invokes `#%app` with explicit instantiation types. In this manner, a surface language with inferred instantiation rewrites to one with explicit instantiation. The second `#%app` clause resembles the first except it does not use the expected type. The third instantiates the polymorphic function type (line 39) and then checks function arguments as in figure 11.

Algebraic datatypes Figure 19's `define-type` macro enables sum-of-products algebraic datatype definitions in MLISH; it expands to a series of definitions (gray box): a type constructor (lines 3-5), where the `#:extra` argument communicates information about the type to other type rules, e.g., to check `match` clause completeness; Racket `struct`s (line 6) implementing runtime constructors; and type rules (lines 7-11) that leverage `#%app` to instantiate polymorphic constructors.

Pattern matching In figure 19's `match`, (one or more) clauses follow e_m, matching its possible variants. The rule uses "extra" information in the type representation to check exhaustiveness of clauses (lines 14-17). Otherwise `match` expands to a conditional that extracts components of e_m using appropriate predicates and accessors (that are also derived from the "extra" information). Here is an MLISH example:

```
#lang mlish
(define-type (Tree X)
  (leaf [val : X]) (node [l : (Tree X)] [r : (Tree X)]))
(define (sum-tr [t : (Tree Int)] → Int)
  (match t with [node l r -> (+ (sum-tr l) (sum-tr r))]))
; TYERR: match: not enough clauses, missing leaf
```

Type classes require intertwined typechecking and program rewriting. Figure 20 sketches an implementation, demonstrating how TURNSTILE's combined rules naturally accommodate such a system. MLISH supports type classes with basic features such as subclassing (unsupported features include multi-parameter type classes and overlapping instances). For simplicity, this paper explains single-operation type classes, though MLISH supports general type classes. The `define-tc` form shows that two definitions implement a type class: a macro for the type class itself (line 2) that expands to its generic operation and type, and a type rule for that operation (lines 3-7) that looks up a concrete operation (line 5) based on the generic name and the concrete types of its usage. MLISH type classes reuse the compile-time macro environment for lookups, where a concrete operation's name, installed by `define-instance` (lines 8-14), is a mangling of the generic name and specific concrete types.

Consequently, functions utilizing generic operations (not shown) have a typeclass component in their type (the `=>` constructor on line 18) and these functions implicitly have an extra concrete operation argument. The `#%app` rule implicitly inserts this argument by: extracting the generic operation of the type class (line 21); looking up the concrete operation

test description	core langs (§ 5)	MLISH (§ 6)
coverage	4313	2467
RW OCaml [32]	———	610
Benchmarks Game [1]	———	852
Okasaki [35]	———	2014
Other examples (e.g., nqueens)	———	559
total (LoC, incl. comments)	4313	6502

Table 3. Testing TURNSTILE-created languages

based on instantiation types for the function (lines 22-23); and adding this operation to the application (line 25).

7. Creating a Test Suite

Sections 5 and 6 show that our approach accommodates a variety of typed languages. This section explains how we validate these languages with a large test suite that includes a variety of real-world programs [1, 2, 32, 35]. To implement this test suite, we created a type system unit-testing framework that considers both typechecking successes and failures, including error messages. It also enables writing tests in a language's surface syntax rather than an internal AST structure. The following defines a function f, and tests the type of f and both a successful and failing application of f:

```
#lang mlish    (require typechecker-tester)
(define (f [x : Int] → Int) x)
(check-type f : (→ Int Int))
(check-type (f 1) : Int => 1)
(typecheck-fail (f 1 2) #:msg ''Wrong number of args'')
```

Table 3 summarizes our test suite, which includes both "coverage" tests checking general functionality and corner cases, and real-world examples. For the latter, Real World OCaml [32] supplied functional tests while the Benchmarks Game [1] consisted of more imperative tests. Okasaki's data structures tested the limits of our type system. For example, in discussing polymorphic recursion (ch. 10), he writes:

> "We will often present code as if SML supports [polymor. recursion]. This code will not be executable but will be easier to read."

We were able to add polymorphic recursion to MLISH, by leveraging recursive definition forms in the host, and to implement the data structures in question, demonstrating both the ability to implement tricky type system features with TURNSTILE, and the ease with which one can do so.

8. Related Work

Researchers have developed many frameworks for writing *type system extensions* [3, 5, 12, 14, 31, 36, 48]. These systems, however, are designed to extend their host languages and typically defer to the host's type system to compute type checking. As such, any new rules must be expressible with the underlying type system. In contrast, our approach does not require commitment to a particular type system but a programmer may have to implement more type rules. Our approach attempts to simplify the latter as much as possible.

Typed Racket [41] pioneered the idea of implementing a typed language with syntax extensions of an untyped language. While our paper shares this approach at a high-level,

the novelty of our work stems not from the use of macros to implement typed languages but the manner in which we do so. Specifically, Typed Racket builds a type checker from scratch and calls out to it after macro expansion. Since it does not reuse Racket's implementation for type checking, Typed Racket is considered a "sister" language [42] to Racket rather than a Racket-embedded language. Further, Typed Racket's type checker is not easily reused to help implement other typed languages, as users attest [23, 27]. We conjecture that reimplementing Typed Racket in our framework would improve its reusability. The availability of type information during macro expansion may also enable new type features such as generic operations, a feature that Typed Racket has struggled to implement.

The ***TinkerType*** [30] system also separates type rules and type operations into reusable components, which resembles our modularization of type systems into macro components. Though the framework stitches together raw string components, and is designed for modeling and typesetting typed language calculi rather than creating practical languages, our approach may benefit from some of the consistency checks that TinkerType performs when combining components. We plan to explore this in future work; for now, the time-honored task of proving type soundness remains for human experts.

9. Conclusions and Future Work

We present a novel use of macros to create practical, performant typed embedded languages. Our approach is not constrained to a particular type system, yet programmers do not have to implement a system from scratch because they can reuse the infrastructure of a macro system. To this end we introduce TURNSTILE, a metalanguage that enables creating typed languages by writing combined type and rewriting rules in an intuitive mathematical syntax. We conjecture that many typed language implementers will benefit from our approach, as non-experts may reduce the burden of language creation, while language researchers may rapidly iterate and experiment with new type features and combinations of features without implementing entire languages.

As next steps, we plan to further validate our approach with the implementation of more languages, and to extend our approach by exploring its application to more complex static analyses. In addition, while this paper presents synergies between type checking and one particular Lisp-style macro system we plan to explore whether the same approach to implementing typed embedded languages is compatible with other syntax extension systems. Finally, we plan to investigate whether the connections between type checking and macro processing that we have described might inform the *future* design of both type systems and macro systems. The fact that many type systems already intertwine type checking and program rewrites [38, 39, 46] suggests that perhaps they should be computed together *by default*. We plan to explore the ramifications of such integration.

References

[1] The computer language benchmarks game. URL `http://benchmarksgame.alioth.debian.org/`.

[2] H. Abelson and G. J. Sussman. *Structure and Interpretation of Computer Programs*. Cambridge University Press, 1979.

[3] C. Andreae, J. Noble, S. Markstrum, and T. Millstein. A framework for implementing pluggable type systems. In *Proceedings of the 21st Annual ACM SIGPLAN Conference on Object-oriented Programming Systems, Languages, and Applications*, pages 57–74, 2006.

[4] P. Bagwell. DSLs - A powerful Scala feature, 2009. URL `http://www.scala-lang.org/old/node/1403`.

[5] G. Bracha. Pluggable type systems. In *OOPSLA Workshop on Revival of Dynamic Languages*, 2004.

[6] N. G. D. Bruijn. Lambda calculus notation with nameless dummies, a tool for automatic formula manipulation, with application to the Church-Rosser theorem. *INDAG. MATH*, 34:381–392, 1972.

[7] E. Burmako. Scala macros: Let our powers combine!: On how rich syntax and static types work with metaprogramming. In *Proceedings of the 4th Workshop on Scala*, 2013.

[8] M. Butterick. Pollen: the book is a program, 2013. URL `https://github.com/mbutterick/pollen`.

[9] N. Cameron. Sets of scopes macro hygiene in Rust, 2015. URL `http://www.ncameron.org/blog/sets-of-scopes-macro-hygiene/`.

[10] R. Culpepper and M. Felleisen. Fortifying macros. In *Proceeding of the 15th ACM SIGPLAN International Conference on Functional Programming*, pages 235–246, 2010.

[11] T. Disney. *Hygienic Macros for JavaScript*. PhD thesis, University of California Santa Cruz, 2015.

[12] S. Erdweg, T. Rendel, C. Kästner, and K. Ostermann. Sugarj: Library-based syntactic language extensibility. In *Proceedings of the ACM SIGPLAN International Conference on Object Oriented Programming Systems Languages and Applications*, pages 391–406, 2011.

[13] M. Felleisen, R. B. Findler, M. Flatt, S. Krishnamurthi, E. Barzilay, J. McCarthy, and S. Tobin-Hochstadt. The Racket Manifesto. In *1st Summit on Advances in Programming Languages (SNAPL 2015)*, pages 113–128, 2015.

[14] D. Fisher and O. Shivers. Static analysis for syntax objects. In *Proceedings of the 11th ACM SIGPLAN International Conference on Functional Programming*, pages 111–121, 2006.

[15] M. Flatt. Binding as sets of scopes. In *Proceedings of the 43rd ACM SIGPLAN-SIGACT Symposium on Programming Languages*, pages 705–717, 2016.

[16] M. Flatt and PLT. Reference: Racket. Technical Report PLT-TR-2010-1, PLT Design Inc., 2010. `http://racket-lang.org/tr1/`.

[17] M. Flatt, R. Culpepper, D. Darais, and R. B. Findler. Macros that work together. *Journal of Functional Programming*, 22 (2):181–216, 2012.

[18] M. Fowler and R. Parsons. *Domain-Specific Languages*. Addison-Wesley, 2010.

[19] S. Freeman, T. Mackinnon, N. Pryce, and J. Walnes. Mock roles, not objects. In *Companion to the 19th Annual ACM SIGPLAN Conference on Object-oriented Programming Systems, Languages, and Applications*, pages 236–246, 2004.

[20] M. Gasbichler and M. Sperber. Integrating user-level threads with processes in scsh. *Higher Order Symbol. Comput.*, 18 (3-4):327–354, 2005.

[21] P. Graham. *On Lisp*. Prentice Hall, 1993.

[22] D. Herman, L. Wagner, and A. Zakai. asm.js working draft, 2014. URL `http://asmjs.org/spec/latest/`.

[23] K. Hinsen. Defining a typed language. Racket users mailing list, 2014. URL `https://groups.google.com/forum/#!topic/racket-users/5LWRHCc6Q6Y`.

[24] P. Hudak. Building domain-specific embedded languages. *ACM Comput. Surv.*, 28(4es), 1996.

[25] P. Hudak. Modular domain specific languages and tools. In *Proceedings of the 5th International Conference on Software Reuse*, pages 134–142, 1998.

[26] S. Kamin and D. Hyatt. A special-purpose language for picture-drawing. In *Proceedings of the USENIX Conference on Domain-Specific Languages*, pages 297–310, 1997.

[27] A. King. Creating a language that extends typed/racket. Racket users mailing list, 2015. URL `https://groups.google.com/forum/#!topic/racket-users/KFx1P7rUK0c`.

[28] G. Kossakowski, N. Amin, T. Rompf, and M. Odersky. JavaScript as an embedded DSL. In *European Conference on Object-Oriented Programming*, pages 409–434. Springer Berlin Heidelberg, 2012.

[29] S. Krishnamurthi. *Linguistic Reuse*. PhD thesis, Rice University, 2001.

[30] M. Y. Levin and B. C. Pierce. TinkerType: A language for playing with formal systems. *J. Funct. Program.*, 13(2):295–316, 2003.

[31] F. Lorenzen and S. Erdweg. Sound type-dependent syntactic language extension. In *Proceedings of the 43rd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, pages 204–216, 2016.

[32] Y. Minsky, A. Madhavapeddy, and J. Hickey. *Real World OCaml*. O'Reilly Media, 2013.

[33] S. Moore, C. Dimoulas, D. King, and S. Chong. Shill: A secure shell scripting language. In *Proceedings of the 11th USENIX Conference on Operating Systems Design and Implementation*, pages 183–199, 2014.

[34] F. Nielson and H. R. Nielson. *Correct System Design: Recent Insights and Advances*, chapter Type and Effect Systems, pages 114–136. Springer Berlin Heidelberg, 1999.

[35] C. Okasaki. *Purely Functional Data Structures*. Cambridge University Press, 1999.

[36] C. Omar, D. Kurilova, L. Nistor, B. Chung, A. Potanin, and J. Aldrich. Safely composable type-specific languages. In *European Conference on Object-Oriented Programming*, pages 105–130. Springer Berlin Heidelberg, 2014.

[37] B. C. Pierce. *Types and Programming Languages*. MIT Press, 2002.

[38] B. C. Pierce and D. N. Turner. Local type inference. In *Proceedings of the 25th ACM SIGPLAN–SIGACT Symposium on Principles of Programming Languages*, pages 252–265, 1998.

[39] J. G. Siek, M. M. Vitousek, M. Cimini, and J. T. Boyland. Refined criteria for gradual typing. In *1st Summit on Advances in Programming Languages*, pages 274–293, 2015.

[40] T. S. Strickland, B. M. Ren, and J. S. Foster. Contracts for domain-specific languages in Ruby. In *Proceedings of the 10th ACM Symposium on Dynamic Languages*, pages 23–34, 2014.

[41] S. Tobin-Hochstadt and M. Felleisen. The design and implementation of Typed Scheme. In *Proceedings of the 35th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, pages 395–406, 2008.

[42] S. Tobin-Hochstadt, V. St-Amour, R. Culpepper, M. Flatt, and M. Felleisen. Languages as libraries. In *Proceedings of the 32nd ACM SIGPLAN Conference on Programming Language Design and Implementation*, pages 132–141, 2011.

[43] E. Torlak and R. Bodik. A lightweight symbolic virtual machine for solver-aided host languages. In *Proceedings of the 35th ACM SIGPLAN Conference on Programming Language Design and Implementation*, pages 530–541, 2014.

[44] B. Valiron, N. J. Ross, P. Selinger, D. S. Alexander, and J. M. Smith. Programming the quantum future. *Communications of the ACM*, 58(8):52–61, 2015.

[45] T. Veldhuizen. Using C++ template metaprograms. *C++ Report*, 7(4):36–43, May 1995.

[46] P. Wadler and S. Blott. How to make ad-hoc polymorphism less ad-hoc. In *Proceedings of the 16th ACM Symposium on Principles of Programming Languages*, pages 60–76, 1989.

[47] J. Yallop and L. White. Modular macros. In *Proceedings of the OCaml Users and Developers Workshop*, 2015.

[48] D. Zook, S. S. Huang, and Y. Smaragdakis. Generating AspectJ programs with Meta-AspectJ. In *Proceedings of the Third International Conference on Generative Programming and Component Engineering*, pages 1–18, 2004.

A.   Syntax Conventions in the Paper

This appendix summarizes terminology and syntax conventions from paper. This paper utilizes pseudocode whenever it aids readability.

- syntax object: An s-expression of symbols, but infused with context information. This is Racket's data representation of programs, replacing the traditional AST.

- `[square brackets]`: Square brackets in Racket are semantically equivalent to regular parentheses and are used for readability purposes, typically to denote binding pairs or language forms with multiple clauses.

- `a syntax pattern`: Syntax patterns deconstruct syntax objects and bind pattern variables. This paper highlights them in yellow. Pattern variables may additionally be tagged with a superscript "syntax class" [10], e.g.,

`type`, which specifies additional constraints that must be satisfied in order for the matching to succeed.

- `a syntax template`: Dual to syntax patterns, syntax templates construct syntax objects, possibly referencing pattern variables. This paper highlights them in gray. Racket code more typically uses the `#'` constructor.

- **syntactic literals**: Bolded, Roman-font identifiers in syntax patterns represent literal symbols that must be matched exactly (as opposed to pattern variables). Examples include the colon to separate a function's parameters and types: $(\lambda \ ([x : \tau]) \ x)$ and the **with** in $(\texttt{match e } \textbf{with} \ \rule{1.5em}{0.4pt})$.

- $\overline{e}$: A version of typed-term e where types are erased. We also use the overline to disambiguate Racket's forms (e.g., $\overline{\lambda}$) from the typechecked version of the form we wish to define in a typed surface language (e.g., λ). Typically, the overline is used in the output syntax object of macro definitions.

- ; a comment: Comments in Racket begin with ";" and extend to the end of the line. This paper uses a gray Roman font face to distinguish comments.

- $\overrightarrow{\tau=}$: The over-arrow notation represents a function lifted to consume compound, e.g. list, inputs. For example ($\overrightarrow{\tau=}$ $\tau s_1 \ \tau s_2$) is equivalent to repeatedly applying $\tau=$ to each pair of types in τs_1 and τs_2.

- $\boxed{\text{MODULENAME}}$: Module names are boxed, use the small-caps style, and are typically located in a figure's upper-right corner. Module names are also language names, e.g., STLC, and typically (but not always) correspond to file names.

- $\boxed{\text{param}}$: Overridable hooks in TURNSTILE. Used to implement type operations, e.g., $\tau=$, in an extensible manner.

- $\rule{3em}{0.4pt}$: This long dash indicates elided code. We also sometimes use a four-dot ellipses for this purpose (because the three-dot ellipses is a valid syntax pattern).

B. TURNSTILE vs Racket syntax

This appendix gives an idea of how TURNSTILE's syntax matches up with equivalent Racket macro code.

Premises:

TURNSTILE	Racket
$\vdash e \gg \overline{e} \Rightarrow \tau$	`#:with` $[\overline{e}\ \tau]$ `(comp+erase-`τ` e)`

$\vdash e \gg \overline{e} \Leftarrow \tau_{\text{expected}}$	`#:with` $[\overline{e}\ \tau]$ `(comp+erase-`τ` (add-expected-`τ` e ` τ_{expected}`))`
	
	`#:fail-unless (`$\boxed{\tau \sqsubseteq}$ τ τ_{expected}`)`
	`                (typecheck-fail-msg ` τ_{expected} τ ` e)`

$[x \gg \overline{x} : \tau_x] \vdash e \gg \overline{e} \Rightarrow \tau$	`#:with` $[(\overline{x})\ \overline{e}\ \tau]$ `(comp+erase-`τ`/ctx e (`$[x\ \tau_x]$`))`

$[x \gg \overline{x} : \tau_x] \vdash e \gg \overline{e} \Leftarrow \tau_{\text{expected}}$	`#:with` $[(\overline{x})\ \overline{e}\ \tau]$ `(comp+erase-`τ`/ctx (add-expected-`τ` e ` τ_{expected}`) (`$[x\ \tau_x]$`))`
	
	`#:fail-unless (`$\boxed{\tau \sqsubseteq}$ τ τ_{expected}`)`
	`                (typecheck-fail-msg ` τ_{expected} τ ` e)`

Conclusion: (separated into input patterns and output syntax)

TURNSTILE		Racket	
Input Components	Output Components	Input Components	Output Components
$e \gg$	$\vdash \overline{e} \Rightarrow \tau$	`e`	`(add-`τ $\overline{e}$ τ`)`

$e \Leftarrow \tau_{\text{expected}} \gg$	$\vdash \overline{e}$	`e`
		`#:when (has-expected-`τ`? e)`
		`#:with ` τ_{expected} `(get-expected-`τ` e)`
			`(add-`τ $\overline{e}$ τ_{expected}`)`